

# Profiling with TAU - Basics

Dr Chennakesava Kadapa

Swansea Academy of Advanced Computing (SA2C), Swansea University, UK.

Email: [c.kadapa@swansea.ac.uk](mailto:c.kadapa@swansea.ac.uk)

## **General workflow**

- Preparation
- Measurement
- Analysis
- Examination
- Optimisation

## **Performance analysis steps**

- Program instrumentation
- Summary measurement collection
- Summary analysis report examination

## TAU + MPI

- **Some of the features (unwinding) are not supported by Intel MPI.**
- PAPI counters do not work with Intel MPI.
- The MPI library used with TAU is OpenMPI (3.1.1).

## Modules on Sunbird

- To use TAU, we need to load the appropriate modules first.
- TAU is available for use with Intel and GNU compilers on *Sunbird*.
- The default compiler is GNU 8.1.0
  - `module purge`
  - `module load tau/2.28.1`
- To use TAU with Intel/2018/2 compiler, we need to load the compiler first.
  - `module purge`
  - `module load compiler/gnu/8/1.0`
  - `module load tau/2.28.1`

# Instrumentation methods

TAU provides a list of instrumentation methods. Among them, the following three are popular.

- Dynamic instrumentation
- Source instrumentation
- Selective instrumentation

# Dynamic instrumentation

# Dynamic instrumentation

- The most straightforward way of instrumentation.
- Does not require modifications to the binary. (Build the executable using `mpicxx .`)
- It makes use of statistical sampling to estimate the percentage of execution times taken by each function.
- Add `tau_exec` before the name of the binary.
  - `mpirun -n 10 ./a.out`
  - `mpirun -n 10 tau_exec ./a.out`
  - `mpirun -n 10 tau_exec -T mpi,papi -ompt ./a.out`
- Unfortunately, this method doesn't support profiling user-defined routines but only those from the MPI library.



## Dynamic instrumentation - Example

- Use the wave2d.cpp code from <https://github.com/chennachaos/TAUhandson>
- **Compile:** `mpicxx wave2d.cpp -o wave2d`
- **Run:** `mpirun -n 12 tau_exec ./wave2d 500 500 3 4 5`
- Once the run is successfully completed, we will get 12 files in `profile.a.b.c` format. `a` is the rank of processor, `b` is the context number and `c` is the thread number.

## Visualisation of profiling data

- **pprof**: Text based visualisation.
  - Check with `pprof --help` for the list of options
- **paraprof**: GUI based visualisation.

# Dynamic instrumentation - Example - Output with pprof

- Visualisation of profiling data using `pprof`.

```
[s.engkadac@sl1 ex01-without-options-without-env]$ pprof -s -n 10 -p  
Reading Profile files in profile.*
```

FUNCTION SUMMARY (total):

%Time	Exclusive msec	Inclusive total msec	#Call	#Subrs	Inclusive usec/call	Name
100.0	24039	39002	16	14496	2437603	.TAU application
33.8	13174	13174	16	0	823378	MPI_Init()
2.4	944	944	16	0	58988	MPI_Finalize()
1.8	712	712	6400	0	111	MPI_Send()
0.3	98.1	98.1	1600	0	61	MPI_Waitall()
0.1	29.6	29.6	16	0	1853	MPI_Bcast()
0.0	5.08	5.08	6400	0	1	MPI_Irecv()
0.0	0.021	0.021	32	0	1	MPI_Comm_rank()
0.0	0.002	0.002	16	0	0	MPI_Comm_size()


FUNCTION SUMMARY (mean):

%Time	Exclusive msec	Inclusive total msec	#Call	#Subrs	Inclusive usec/call	Name
100.0	1502	2438	1	906	2437603	.TAU application
33.8	823	823	1	0	823378	MPI_Init()
2.4	59.0	59.0	1	0	58988	MPI_Finalize()
1.8	44.5	44.5	400	0	111	MPI_Send()
0.3	6.13	6.13	100	0	61	MPI_Waitall()
0.1	1.85	1.85	1	0	1853	MPI_Bcast()
0.0	0.318	0.318	400	0	1	MPI_Irecv()
0.0	0.001	0.001	2	0	1	MPI_Comm_rank()
0.0	0.000	0.000	1	0	0	MPI_Comm_size()

# Dynamic instrumentation - Example - Output with paraprof


- To visualise the profiling data, we prefer `paraprof`.
- Call `paraprof` from the command line. If you are connecting remotely, then you need to connect with X server activated; `ssh -X sunbird`.

Metric: TIME  
Value: Exclusive



# Dynamic instrumentation - Example - Output with paraprof

Metric: TIME  
Value: Exclusive  
Units: seconds



# Dynamic instrumentation - Example - Output

Name: MPI\_Init()  
Metric Name: TIME  
Value: Exclusive  
Units: seconds



# Compiler-based instrumentation

- Compiler instrumentation method sits in between *dynamic* and *source* instrumentation.
- It requires compilation of the source code with additional arguments.
- It cannot provide information about finer constructs such as loops.
- To use this method, we should add `-tau-options=-optCompInst` to `tau_cxx.sh`.
- There is not much we can understand about the performance of code using this method.
- Recommended practice is to follow the *source* instrumentation which is fine-grain.



## TAU's environment variables

Variable	Description
TAU_PROFILE	Set to 1 to have TAU profile your code
TAU_TRACE	Set to 1 to have TAU trace your code
TAU_SAMPLING	Default value is zero (i.e., Sampling is OFF by default). When set to 1, TAU will collect additional profile or trace information via periodic sampling at runtime

## TAU's environment variables (cont'd)

Variable	Description
TAU_CALLPATH	TAU will generate call-path data when this is set to 1
TAU_CALLPATH_DEPTH	Sets the depth of the callpath profiling
TAU_VERBOSE	When set TAU will print out information about its configuration when a running a instrumented application

More details at <https://www.cs.uoregon.edu/research/tau/docs/newguide/bk05apa.html>

# Source instrumentation

- Replace `mpicc`, `mpicxx` and `mpif90` with `taucc`, `taucxx` and `tauf90`.
- Add `-dynamic` flag to the linking phase.
- Use files from `ex4`

## Commonly encountered errors

- profile.x.x.x files are not generated.

- Run time error

```
wave2d: ../../scorep-5.0/./build-  
mpi/../../src/measurement/paradigm/mpi/scorep_ipc_mpi.c:230:  
SCOREP_IpcGroup_GetRank: Assertion 'SCOREP_Status_IsMppInitialized()'   
failed.
```

- Turns out that the error is due to the incorrect value of `TAU_MAKEFILE` .

```
TAU_MAKEFILE=/apps/tools/tau/2.28.1/e17/AVX512/gnu-8.1/openmpi-  
3.1/x86_64/lib/Makefile.tau-papi-mpi-pdt-scorep
```

- Need to be extra careful when setting the `TAU_MAKEFILE` variable

- Works correctly after using

```
TAU_MAKEFILE=/apps/tools/tau/2.28.1/e17/AVX512/gnu-8.1/openmpi-  
3.1/x86_64/lib/Makefile.tau-papi-mpi-pdt
```

- `TAU_PROFILE`, `TAU_CALLPATH` and `TAU_SAMPLING` are set to 1.
- `-g` flag is not used.

# Source instrumentation - Example - Output with pprof

- Visualisation of profiling data using `pprof`.
- Notice the change in the output. We now see the user-defined functions in the profiling data.

```
[s.engkadac@sl1 ex04-with-taucompilers]$ pprof -s -n 10 -p  
Reading Profile files in profile.*
```


```
FUNCTION SUMMARY (total):
```

%Time	Exclusive msec	Inclusive total msec	#Call	#Subrs	Inclusive Name usec/call
100.0	165	112075	16	16	7004669 .TAU application
99.9	40.8	111910	16	128	6994358 .TAU application => main
99.9	40.8	111910	16	128	6994358 main
93.9	3.81	105287	16	3200	6580467 Grid::doIterations(int)
93.9	3.81	105287	16	3200	6580467 main => Grid::doIterations(int)
92.0	0.000	103140	3438	0	30000 Grid::doIterations(int) => Grid::doOneIteration()
92.0	103140	103140	3438	0	30000 Grid::doIterations(int) => Grid::doOneIteration() => [CONTEXT] Grid::doOneIteration()
92.0	0.000	103140	3438	0	30000 [CONTEXT] Grid::doOneIteration()
92.0	103140	103140	3438	0	30000 [SAMPLE] Grid::doOneIteration()
91.8	102905	102905	1600	0	64316 Grid::doIterations(int) => Grid::doOneIteration()
132.2	130366	148151	34524	28864	4291 -others-

# Source instrumentation - Example - Output with paraprof

- To visualise the profiling data, we prefer `paraprof`.
- Call `paraprof` from the command line. If you are connecting remotely, then you need to connect with X server activated; `ssh -X sunbird`.

Metric: TIME  
Value: Exclusive





# Source instrumentation - Example - Output with paraprof



# Source instrumentation - Example - Output

```
Name: Grid::doIterations(int)
[ {/home/s.engkadam/profiling/wave2d/tau/ex04-with-taucompilers/wave2d.cpp} {50,0} ] =>
Grid::doOneIteration() [ {/home/s.engkadam/profiling/wave2d/tau/ex04-with-taucompilers/wave2d.cpp}
{94,0} ] => [CONTEXT] Grid::doOneIteration()
[ {/home/s.engkadam/profiling/wave2d/tau/ex04-with-taucompilers/wave2d.cpp} {94,0} ] =>
[SUMMARY] Grid::doOneIteration()
[ {/home/s.engkadam/profiling/wave2d/tau/ex04-with-taucompilers/wave2d.cpp} ] => [SAMPLE]
Grid::doOneIteration() [ {/home/s.engkadam/profiling/wave2d/tau/ex04-with-taucompilers/wave2d.cpp}
{101} ]
Metric Name: TIME
Value: Exclusive
Units: seconds
```



# Tracing

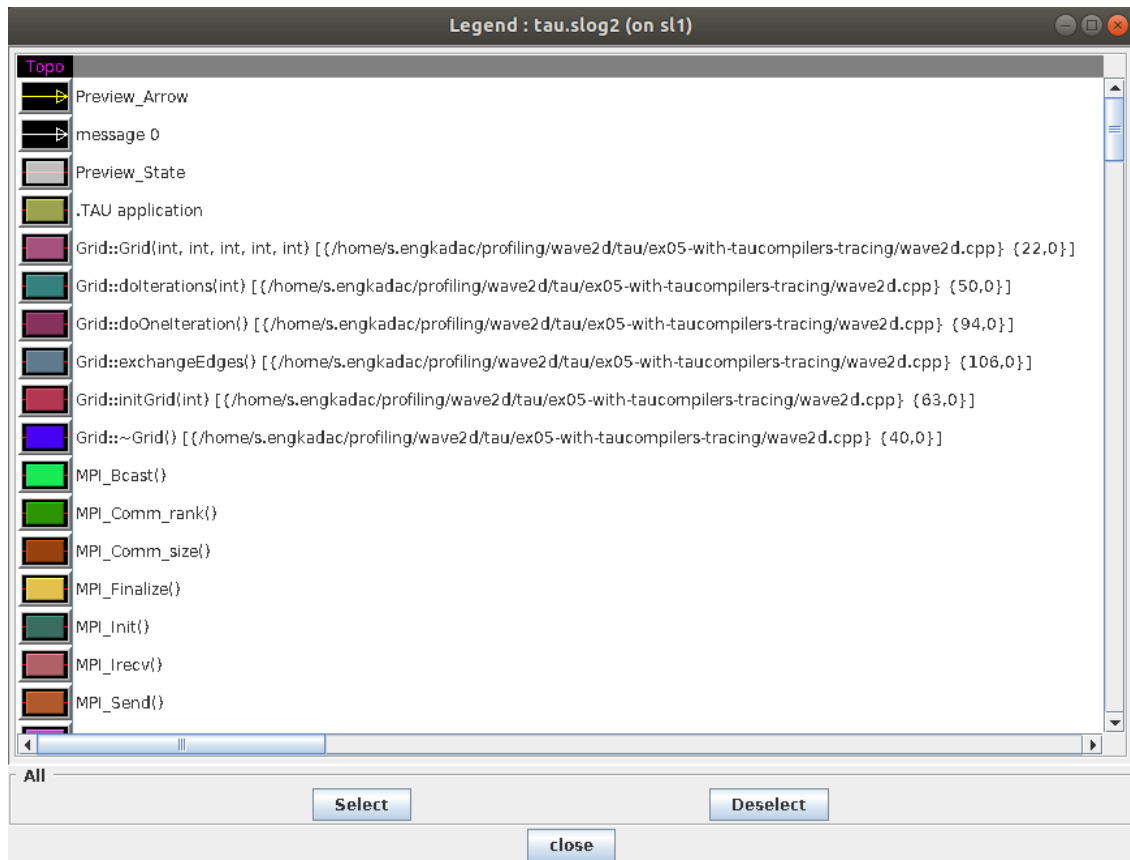
## Tracing with TAU

- TAU can automatically instrument the code to do tracing without user intervention (unlike MPE, which requires manual insertion of tracing code).
- This tracing functionality depends on PDT, its C/C++ parser in particular.
- To enable this, we need to the TAU Makefile with the *PDT* option.
- **Note:** tracing consumes an enormous amount of disk space. So, use it with caution.

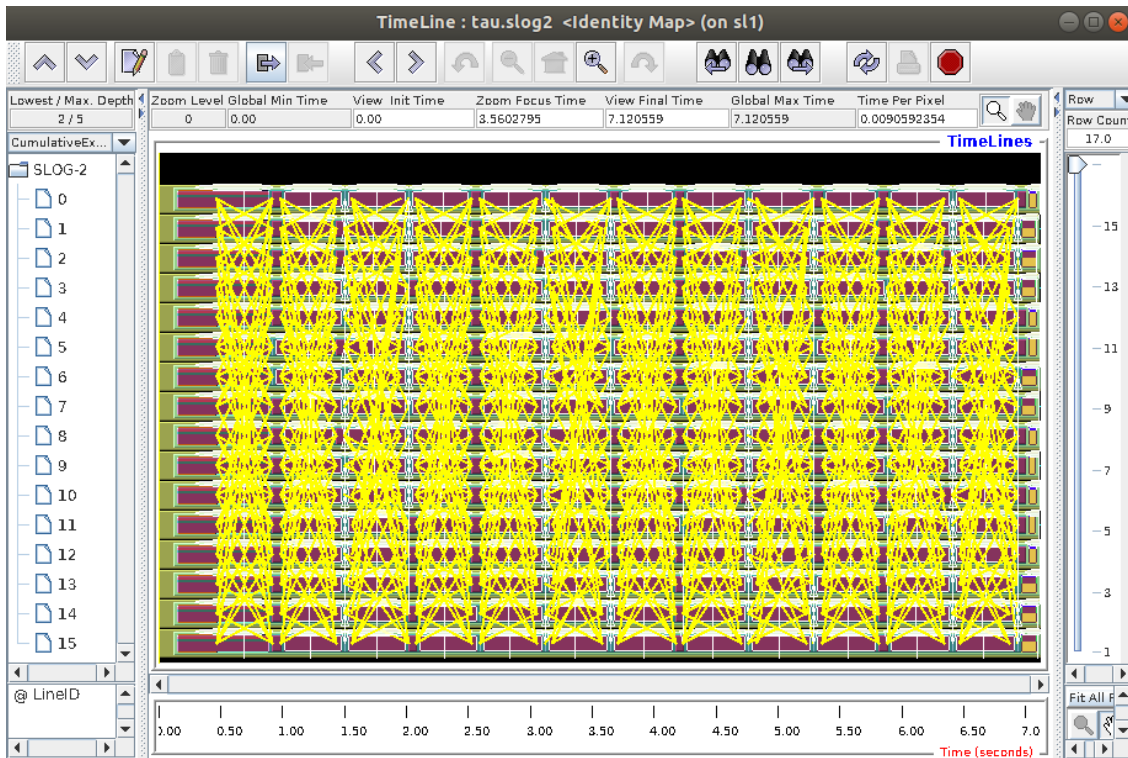
# Tracing with TAU (ex5)

- Tracing is disabled by default. To enable it
  - `export TAU_TRACE=1`
- Enable tracing disables profiling. That is, no `profile.*` files will be generated.
- The execution of the program will generate `.edf` and `.trc` files. These files must be merged first before visualisation.
  - `tau_treemerge.pl`
  - `tau2slog2 tau.trc tau.edf -o tau.slog2`
- Tracing data is dumped into `tau.slog2` which can be visualised with *Jumpshot*. (TAU doesn't have a visual tracer).

# Visualising tracing data with Jumpshot



# Visualising tracing data with Jumpshot



- We can zoom into the graphs by drawing a window around the portion of interest.

